

YIWEI WANG

Phone: +1 4256255851 ◊ Email: wangyw.evan@gmail.com ◊ Website: <https://wangywust.github.io/>

WORK

Assistant Professor at University of California, Merced	Jan, 2025 -
Research and teaching on natural language processing and vision-language models.	
Postdoc at University of California, Los Angeles	Sep. 2023 - Dec. 2024
In the UCLA NLP group, advised by Prof. Kai-Wei Chang and Prof. Nanyun Peng.	
Applied Scientist in Amazon	Sep. 2022 - Sep. 2023

EDUCATION

Doctor of Philosophy in Computer Science	Jul. 2019 - Mar. 2023
National University of Singapore	
Master of Philosophy in Electronic and Computer Engineering	Aug. 2017 - Jun. 2019
Hong Kong University of Science and Technology	
Bachelor of Science in Information Engineering	Aug. 2013 - Jun. 2017
Southeast University, GPA: 3.92/4.0, Rank: 5/218 (Top 3%)	

RESEARCH PUBLICATION

-
- [1] Yiwei Wang, Muhao Chen, Nanyun Peng, Kai-Wei Chang. Vulnerability of Large Language Models to Output Prefix Jailbreaks: Impact of Positions on Safety, *NAACL 2025*.
 - [2] Chen Wang, Yiwei Wang, Yujun Cai, Bryan Hooi. Tricking Retrievers with Influential Tokens: An Efficient Black-Box Corpus Poisoning Attack, *NAACL 2025*.
 - [3] Baolong Bi, Shenghua Liu, Yiwei Wang, Lingrui Mei, Xueqi Cheng. Is Factuality Enhancement a Free Lunch For LLMs? Better Factuality Can Lead to Worse Context-Faithfulness, *ICLR 2025*.
 - [4] Zhicheng Yang, Yiwei Wang, Yinya Huang, Zhijiang Guo, Wei Shi, Xiongwei Han, Liang Feng, Linqi Song, Xiaodan Liang, Jing Tang. OptiBench meets ReSocratic: Measure and improve LLMs for optimization modeling, *ICLR 2025*.
 - [5] Shaochen Zhong, Yifan Lu, Lize Shao, Bhargav Bhushanam, Xiaocong Du, Louis Feng, Yixin Wan, Yucheng Shi, Daochen Zha, Yiwei Wang, others. MQuAKE-Remastered: Multi-Hop Knowledge Editing Can Only Be Advanced With Reliable Evaluations, *ICLR 2025*.
 - [6] Lingrui Mei, Shenghua Liu, Yiwei Wang, Baolong Bi, Jiayi Mao, Xueqi Cheng. "Not Aligned" is Not "Malicious": Being Careful about Hallucinations of Large Language Models' Jailbreak, *COLING 2025*.
 - [7] Cheng Wang, Yiwei Wang, Bryan Hooi, Yujun Cai, Nanyun Peng, Kai-Wei Chang. Con-ReCall: Detecting Pre-training Data in LLMs via Contrastive Decoding, *COLING 2025*.
 - [8] Yihong Luo, Yuhan Chen, Siya Qiu, Yiwei Wang, Chen Zhang, Yan Zhou, Xiaochun Cao, Jing Tang. Fast Graph Sharpness-Aware Minimization for Enhancing and Accelerating Few-Shot Node Classification, *NeurIPS 2024*.

- [9] Baolong Bi, Shenghua Liu, Yiwei Wang, Lingrui Mei, Hongcheng Gao, Yilong Xu, Xueqi Cheng. Adaptive Token Biased: Knowledge Editing via Biasing Key Entities, *EMNLP 2024*.
- [10] Bingxuan Li, Yiwei Wang, Tao Meng, Nanyun Peng, Kai-Wei Chang. Control Large Language Models via Divide-and-Conquer, *EMNLP 2024*.
- [11] Lingrui Mei, Shenghua Liu, Yiwei Wang, Baolong Bi, Xueqi Chen. SLANG: New Concept Comprehension of Large Language Models, *EMNLP 2024*.
- [12] Zhicheng Yang, Yinya Huang, Jing Xiong, Liang Feng, Xiaodan Liang, Yiwei Wang, Jing Tang. AlignedCoT: Prompting Large Language Models via Native-Speaking Demonstrations, *EMNLP 2024*.
- [13] Silin Meng, Yiwei Wang, Cheng-Fu Yang, Nanyun Peng, Kai-Wei Chang. LLM-A*: Large Language Model Enhanced Incremental Heuristic Search on Path Planning, *EMNLP 2024*.
- [14] Baolong Bi, Shenghua Liu, Yiwei Wang, Lingrui Mei, Xueqi Cheng. Lpnl: Scalable link prediction with large language models, *ACL 2024*.
- [15] Juncheng Liu, Bryan Hooi, Kenji Kawaguchi, Yiwei Wang, Chaosheng Dong, Xiaokui Xiao. Scalable and Effective Implicit Graph Neural Networks on Large Graphs, *ICLR 2024*.
- [16] Yuyao Ge, Shenghua Liu, Baolong Bi, Yiwei Wang, Lingrui Mei, Wenjie Feng, Lizhe Chen, Xueqi Cheng. Can Graph Descriptive Order Affect Solving Graph Problems with LLMs?, *2024*.
- [17] Fei Wang, Wenjie Mo, Yiwei Wang, Wenxuan Zhou, Muhao Chen. A Causal View of Entity Bias in (Large) Language Models, *EMNLP 2023*.
- [18] Yiwei Wang, Yujun Cai, Muhao Chen, Yuxuan Liang, Bryan Hooi. Primacy Effect of ChatGPT, *EMNLP 2023*.
- [19] Yuxuan Liang, Yutong Xia, Songyu Ke, Yiwei Wang, Qingsong Wen, Junbo Zhang, Yu Zheng, Roger Zimmermann. Airformer: Predicting nationwide air quality in china with transformers, *AAAI 2023*.
- [20] Yiwei Wang, Bryan Hooi, Fei Wang, Yujun Cai, Yuxuan Liang, Wenxuan Zhou, Jing Tang, Manjuan Duan, Muhao Chen. How Fragile is Relation Extraction under Entity Replacements?, *CONLL 2023*.
- [21] Yiwei Wang, Bryan Hooi, Yozen Liu, Neil Shah. Graph explicit neural networks: Explicitly encoding graphs for efficient and accurate inference, *WSDM 2023*.
- [22] Juncheng Liu, Zequn Sun, Bryan Hooi, Yiwei Wang, Dayiheng Liu, Baosong Yang, Xiaokui Xiao, Muhao Chen. Dangling-Aware Entity Alignment with Mixed High-Order Proximities, *NAACL 2022*.
- [23] Yiwei Wang, Muhao Chen, Wenxuan Zhou, Yujun Cai, Yuxuan Liang, Bryan Hooi. GraphCache: Message Passing as Caching for Sentence-Level Relation Extraction, *NAACL 2022*.
- [24] Yiwei Wang, Muhao Chen, Wenxuan Zhou, Yujun Cai, Yuxuan Liang, Dayiheng Liu, Baosong Yang, Juncheng Liu, Bryan Hooi. Should We Rely on Entity Mentions for Relation Extraction? Debiasing Relation Extraction with Counterfactual Analysis, *NAACL 2022*.
- [25] Juncheng Liu, Yiwei Wang, Bryan Hooi, Renchi Yang, Xiaokui Xiao. Lscale: latent space clustering-based active learning for node classification, *ECML 2022*.
- [26] Yiwei Wang, Bryan Hooi, Yozen Liu, Tong Zhao, Zhichun Guo, Neil Shah. Flashlight: Scalable link prediction with effective decoders, *LOG 2022*.

- [27] Yuxuan Liang, Kun Ouyang, Yiwei Wang, Zheyi Pan, Yifang Yin, Hongyang Chen, Junbo Zhang, Yu Zheng, David S Rosenblum, Roger Zimmermann. Mixed-order relation-aware recurrent neural networks for spatio-temporal forecasting, *TKDE 2022*.
- [28] Juncheng Liu, Kenji Kawaguchi, Bryan Hooi, Yiwei Wang, Xiaokui Xiao. Eignn: Efficient infinite-depth graph neural networks, *NeurIPS 2021*.
- [29] Yiwei Wang, Yujun Cai, Yuxuan Liang, Henghui Ding, Changhu Wang, Siddharth Bhatia, Bryan Hooi. Adaptive data augmentation on temporal graphs, *NeurIPS 2021*.
- [30] Yujun Cai, Yiwei Wang, Yiheng Zhu, Tat-Jen Cham, Jianfei Cai, Junsong Yuan, Jun Liu, Chuanxia Zheng, Sijie Yan, Henghui Ding, others. A unified 3d human motion synthesis model via conditional variational auto-encoder, *ICCV 2021*.
- [31] Yiwei Wang, Yujun Cai, Yuxuan Liang, Henghui Ding, Changhu Wang, Bryan Hooi. Time-aware neighbor sampling for temporal graph networks, *2022 International Joint Conference on Neural Networks (IJCNN) 2021*.
- [32] Yuxuan Liang, Kun Ouyang, Yiwei Wang, Ye Liu, Junbo Zhang, Yu Zheng, David S Rosenblum. Revisiting convolutional neural networks for citywide crowd flow analytics, *ECML 2021*.
- [33] Yiwei Wang, Wei Wang, Yuxuan Liang, Yujun Cai, Bryan Hooi. Progressive supervision for node classification, *ECML 2021*.
- [34] Yuxuan Liang, Kun Ouyang, Hanshu Yan, Yiwei Wang, Zekun Tong, Roger Zimmermann. Modeling Trajectories with Neural Ordinary Differential Equations, *IJCAI 2021*.
- [35] Yuxuan Liang, Kun Ouyang, Junkai Sun, Yiwei Wang, Junbo Zhang, Yu Zheng, David Rosenblum, Roger Zimmermann. Fine-grained urban flow prediction, *WWW 2021*.
- [36] Yiwei Wang, Wei Wang, Yuxuan Liang, Yujun Cai, Bryan Hooi. Curgraph: Curriculum learning for graph classification, *WWW 2021*.
- [37] Yiwei Wang, Wei Wang, Yuxuan Liang, Yujun Cai, Bryan Hooi. Mixup for node and graph classification, *WWW 2021*.
- [38] Yuxuan Liang, Kun Ouyang, Yiwei Wang, David Samuel Rosenblum. Revisiting convolutional neural networks for urban flow analytics, *2020*.
- [39] Yujun Cai, Lin Huang, Yiwei Wang, Tat-Jen Cham, Jianfei Cai, Junsong Yuan, Jun Liu, Xu Yang, Yiheng Zhu, Xiaohui Shen, others. Learning progressive joint propagation for human motion prediction, *ECCV 2020*.
- [40] Yiwei Wang, Shenghua Liu, Minji Yoon, Hemank Lamba, Wei Wang, Christos Faloutsos, Bryan Hooi. Provably robust node classification via low-pass message passing, *ICDM 2020*.
- [41] Yiwei Wang, Wei Wang, Yujun Cai, Bryan Hooi, Beng Chin Ooi. Detecting implementation bugs in graph convolutional network based node classifiers, *ISSRE 2020*.
- [42] Yiwei Wang, Wei Wang, Yuxuan Liang, Yujun Cai, Juncheng Liu, Bryan Hooi. NodeAug: Semi-supervised node classification with data augmentation, *KDD 2020*.

- [43] Licheng Zhao, Yiwei Wang, Sandeep Kumar, Daniel P Palomar. Optimization algorithms for graph Laplacian estimation via ADMM and MM, *TSP 2019*.
- [44] Yiwei Wang, Mark James Carman, Yuan-Fang Li. Using knowledge graphs to explain entity co-occurrence in Twitter, *CIKM 2017*.

HONORS AND AWARDS

Dean's Graduate Research Excellence Award (Highest honor of PhD students at NUS with 3 to 4 recipients per year) Jul. 2022

SDSC Dissertation Research Fellowship (10 annual quotas for Singapore-based Ph.D. Students) Mar. 2021

Research Achievement Award from School of Computing of National University of Singapore Sep. 2020

Meritorious Winner of the Interdisciplinary Contest in Modeling (ICM) held by COMAP (**7%**) Apr. 2016

National Scholarship from Ministry of Education of the People's Republic of China (**3%**) Oct. 2015

TEACHING

Instructor, New Course on Large Language Models	University of California, Merced, 2025
Teaching Assistant, Big Data Systems for Data Science	National University of Singapore, 2021
Teaching Assistant, Big Data Systems for Data Science	National University of Singapore, 2021
Teaching Assistant, Programming Methodology	National University of Singapore, 2021
Teaching Assistant, Knowledge Discovery and Data Mining and 2021	National University of Singapore, 2020
Teaching Assistant, Parallel Computing	National University of Singapore, 2020
Teaching Assistant, Signal Processing and Communications	HKUST, 2018

SERVICE AS AREA CHAIR

International Conference on Learning Representations (ICLR)	2025
International Conference on Machine Learning (ICML)	2025
Association for Computational Linguistics Rolling Review (ACL ARR)	2025
Conference on Language Modeling (COLM)	2025
International Joint Conference on Artificial Intelligence (IJCAI)	2025
IEEE International Conference on Multimedia and Expo (ICME)	2025
International Joint Conference on Neural Networks (IJCNN)	2025

SERVICE AS CONFERENCE REVIEWER

Conference on Empirical Methods in Natural Language Processing (EMNLP)	2023, 2024
Annual Meeting of the Association for Computational Linguistics (ACL)	2023, 2024
North American Chapter of the Association for Computational Linguistics (NAACL)	2023, 2024
ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD)	2023, 2024
International Conference on Artificial Intelligence and Statistics (AISTATS)	2025
Learning on Graphs Conference (LOG)	2022, 2023, 2024
The Web Conference (WWW)	2023, 2024, 2025
International Conference on Learning Representations (ICLR)	2023, 2024
Conference on Neural Information Processing Systems (NeurIPS)	2022, 2023, 2024
The AAAI Conference on Artificial Intelligence	2021, 2022, 2023, 2024
IEEE/CVF Conference on Computer Vision and Pattern Recognition	2021, 2022, 2023, 2024
International Conference on Computer Vision	2021, 2022, 2023, 2024
European Conference on Computer Vision	2022, 2023, 2024
International Joint Conference on Artificial Intelligence	2022, 2023, 2024
IEEE International Conference on Multimedia and Expo	2022, 2023, 2024
European Conference on Machine Learning and Principles and Practice of Knowledge Discovery in Databases (ECML-PKDD)	2020
International Joint Conference on Neural Networks	2022

SERVICE AS JOURNAL REVIEWER

Neurocomputing
Pattern Recognition
IEEE Transactions on Knowledge and Data Engineering (TKDE)
IET Image Processing
IEEE Transactions on Big Data
Computer Science Review
Cybernetics and Systems
Transactions on Systems, Man and Cybernetics:Systems

ABILITIES AND SKILLS

Languages: Chinese, English (IELTS: **8.0**).

Programming: Advanced with MATLAB, Python. Familiar with C/C++, JAVA.