

YIWEI WANG

Phone: +1 425-625-5851 ◊ Email: wangyw.evan@gmail.com ◊ Website: <https://wangywust.github.io/>

WORK

Assistant Professor at University of California, Merced	Jan, 2025 - Now
Computer Science Department. Research and teaching on natural language processing.	
Postdoc at University of California, Los Angeles	Sep. 2023 - Dec. 2024
In the UCLA NLP group, advised by Prof. Kai-Wei Chang and Prof. Nanyun Peng.	
Applied Scientist in Amazon	Sep. 2022 - Sep. 2023

EDUCATION

Doctor of Philosophy in Computer Science	Jul. 2019 - Mar. 2023
National University of Singapore	
Master of Philosophy in Electronic and Computer Engineering	Aug. 2017 - Jun. 2019
Hong Kong University of Science and Technology	
Bachelor of Science in Information Engineering	Aug. 2013 - Jun. 2017
Southeast University	

HONORS AND AWARDS

NSF Computer and Information Science and Engineering Research Initiation Initiative (CRII)	2025
NVIDIA Academic Grants Program	2025
University of California at Merced Faculty Research Award	2025
Best Paper Finalist in EMNLP 2024	2024
Dean's Graduate Research Excellence Award (Highest honor of PhD students at NUS with 3 to 4 recipients per year)	2022
SDSC Dissertation Research Fellowship (10 annual quotas for Singapore-based Ph.D. Students)	2021
Research Achievement Award from School of Computing of National University of Singapore	2020
National Scholarship from Ministry of Education of the People's Republic of China	2015

RESEARCH PUBLICATION

- [1] Zhenlong Yuan, Xiangyan Qu, Chengxuan Qian, Rui Chen, Jing Tang, Lei Sun, Xiangxiang Chu, Dapeng Zhang, Yiwei Wang, Yujun Cai, Shuo Li. Video-STAR: Reinforcing Open-Vocabulary Action Recognition with Tools, *ICLR 2026*.
- [2] Honghao Fu, Yuan Ouyang, Kai-Wei Chang, Yiwei Wang, Zi Huang, Yujun Cai. ContextNav: Towards Agentic Multimodal In-Context Learning, *ICLR 2026*.

- [3] Haojin Yang, Rui Hu, Zequn Sun, Rui Zhou, Yujun Cai, Yiwei Wang. WavefrontDiffusion: Dynamic Decoding Schedule for Improved Reasoning, *ICLR 2026*.
- [4] Baolong Bi, Shenghua Liu, Yiwei Wang, Yilong Xu, Junfeng Fang, Lingrui Mei, Xueqi Cheng. Parameters vs. Context: Fine-Grained Control of Knowledge Reliance in Language Models, *ICLR 2026*.
- [5] Yike Wu, Yiwei Wang, Yujun Cai. ChainMPQ: Interleaved Text-Image Reasoning Chains for Mitigating Relation Hallucinations, *ICLR 2026*.
- [6] Xingjian Tao, Yiwei Wang, Yujun Cai, Zhicheng Yang, Jing Tang. Are LLMs Really Not Knowledgeable? Mining the Submerged Knowledge in LLMs' Memory, *ICLR 2026*.
- [7] Xuan Shen, Weize Ma, Yufa Zhou, Enhao Tang, Yanyue Xie, Zhengang Li, Yifan Gong, Quanyi Wang, Henghui Ding, Yiwei Wang, Pu Zhao, Jun Lin, Jiuxiang Gu. Fastcar: Cache Attentive Replay for Fast Auto-Regressive Video Generation on the Edge, *ICLR 2026*.
- [8] Zhen Xiong, Yujun Cai, Zhecheng Li, Yiwei Wang. Unveiling the Potential of Diffusion Large Language Model in Controllable Generation, *ICLR 2026*.
- [9] Xuan Shen, Brian Wingenroth, Zichao Wang, Jason Kuen, Wanrong Zhu, Ruiyi Zhang, Yiwei Wang, Lichun Ma, Anqi Liu, Hongfu Liu, Tong Sun, Kevin S. Hawkins, Kate Tasker, G. Caleb Alexander, Jiuxiang Gu. OIDA-QA: A Multimodal Benchmark for Analyzing the Opioid Industry Documents Archive, *AAAI 2026*.
- [10] Yifan Lu, Pengfei Jiao, Xuan Guo, Ziyun Zou, Yiwei Wang, Mengzhou Gao, Huaming Wu, Imran Razzak. Towards Robust Heterogeneous Graph Explanations under Structural Perturbations, *WWW 2026*.
- [11] Juncheng Liu, Chenghao Liu, Gerald Woo, Yiwei Wang, Bryan Hooi, Caiming Xiong, Doyen Sahoo. UniTST: Effectively Modeling Inter-Series and Intra-Series Dependencies for Multivariate Time Series Forecasting, *TMLR 2025*.
- [12] Jiayu Yao, Shenghua Liu, Yiwei Wang, Lingrui Mei, Baolong Bi, Yuyao Ge, Zhecheng Li, Xueqi Cheng. Who is in the Spotlight: The Hidden Bias Undermining Multimodal Retrieval-Augmented Generation, *EMNLP 2025*.
- [13] Zhen Xiong, Yujun Cai, Zhecheng Li, Yiwei Wang. Mapping the Minds of LLMs: A Graph-Based Analysis of Reasoning LLMs, *EMNLP 2025*.
- [14] Shuyang Hao, Yiwei Wang, Bryan Hooi, Jun Liu, Muhan Chen, Zi Huang, Yujun Cai. Making every step effective: Jailbreaking large vision-language models through hierarchical kv equalization, *EMNLP 2025*.
- [15] Haonan Ge, Yiwei Wang, Ming-Hsuan Yang, Yujun Cai. MRFD: Multi-Region Fusion Decoding with Self-Consistency for Mitigating Hallucinations in LVLMs, *EMNLP 2025*.
- [16] Hang Wu, Hongkai Chen, Yujun Cai, Chang Liu, Qingwen Ye, Ming-Hsuan Yang, Yiwei Wang. DiMo-GUI: Advancing Test-time Scaling in GUI Grounding via Modality-Aware Visual Reasoning, *EMNLP 2025*.
- [17] Wen Tao, Yiwei Wang. How to Make Large Language Models Generate 100% Valid Molecules?, *EMNLP 2025*.

- [18] Sifan Li, Yujun Cai, Yiwei Wang. SemVink: Advancing VLMs' Semantic Understanding of Optical Illusions via Visual Global Thinking, *EMNLP 2025*.
- [19] Xingjian Tao, Yiwei Wang, Yujun Cai, Zhicheng Yang, Jing Tang. Understanding GUI Agent Localization Biases through Logit Sharpness, *EMNLP 2025*.
- [20] Zhecheng Li, Guoxian Song, Yujun Cai, Zhen Xiong, Junsong Yuan, Yiwei Wang. Texture or Semantics? Vision-Language Models Get Lost in Font Recognition, *COLM 2025*.
- [21] Zhaochen Wang, Bryan Hooi, Yiwei Wang, Ming-Hsuan Yang, Zi Huang, Yujun Cai. Text Speaks Louder than Vision: ASCII Art Reveals Textual Biases in Vision-Language Models, *COLM 2025*.
- [22] Chunxue Xu, Yiwei Wang, Bryan Hooi, Yujun Cai, Songze Li. How does Watermarking Affect Visual Language Models in Document Understanding?, *COLM 2025*.
- [23] Baolong Bi, Shaohan Huang, Yiwei Wang, Tianchi Yang, Zihan Zhang, Haizhen Huang, Lingrui Mei, Junfeng Fang, Zehao Li, Furu Wei, Weiwei Deng, Feng Sun, Qi Zhang, Shenghua Liu. Context-DPO: Aligning Language Models for Context-Faithfulness, *ACL 2025*.
- [24] Baolong Bi, Shenghua Liu, Lingrui Mei, Yiwei Wang, Junfeng Fang, Pengliang Ji, Xueqi Cheng. Decoding by Contrasting Knowledge: Enhancing Large Language Model Confidence on Edited Facts, *ACL 2025*.
- [25] Yuyao Ge, Shenghua Liu, Baolong Bi, Yiwei Wang, Lingrui Mei, Wenjie Feng, Lizhe Chen, Xueqi Cheng. Can Graph Descriptive Order Affect Solving Graph Problems with LLMs?, *ACL 2025*.
- [26] Zhecheng Li, Yiwei Wang, Bryan Hooi, Yujun Cai, Nanyun Peng, Kai-Wei Chang. DRS: Deep Question Reformulation With Structured Output, *ACL 2025*.
- [27] Bingxuan Li, Yiwei Wang, Jiuxiang Gu, Kai-Wei Chang, Nanyun Peng. METAL: A Multi-Agent Framework for Chart Generation with Test-Time Scaling, *ACL 2025*.
- [28] Zhecheng Li, Yiwei Wang, Bryan Hooi, Yujun Cai, Zhen Xiong, Nanyun Peng, Kai-Wei Chang. Vulnerability of LLMs to Vertically Aligned Text Manipulations, *ACL 2025*.
- [29] Yurong Wu, Fangwen Mu, QiuHong Zhang, Jinjing Zhao, Xinrun Xu, Lingrui Mei, Yang Wu, Lin Shi, Junjie Wang, Zhiming Ding, Yiwei Wang. EvoStealer: Differential Evolution for Prompt Template Stealing Against Text-to-Image Synthesis, *ACL 2025*.
- [30] Rongzhi Zhu, Xiangyu Liu, Zequn Sun, Yiwei Wang, Wei Hu. Mitigating Lost-in-Retrieval Problems in Retrieval Augmented Multi-Hop Question Answering, *ACL 2025*.
- [31] Songlin Zhai, Yuan Meng, Yongrui Chen, Yiwei Wang, Guilin Qi. Peripheral Memory for LLMs: Integration of Sequential Memory Banks with Adaptive Querying, *ICML 2025*.
- [32] Baolong Bi, Shenghua Liu, Yiwei Wang, Lingrui Mei, Xueqi Cheng. Is Factuality Enhancement a Free Lunch For LLMs? Better Factuality Can Lead to Worse Context-Faithfulness, *ICLR 2025*.
- [33] Zhicheng Yang, Yiwei Wang, Yinya Huang, ZhiJiang Guo, Wei Shi, Xiongwei Han, Liang Feng, LinQi Song, Xiaodan Liang, Jing Tang. OptiBench meets ReSocratic: Measure and improve LLMs for optimization modeling, *ICLR 2025*.

- [34] Shaochen Zhong, others. MQuAKE-Remastered: Multi-Hop Knowledge Editing Can Only Be Advanced With Reliable Evaluations, *ICLR 2025*.
- [35] Cheng Wang, Yiwei Wang, Yujun Cai, Bryan Hooi. Tricking Retrievers with Influential Tokens: An Efficient Black-Box Corpus Poisoning Attack, *NAACL 2025*.
- [36] Yiwei Wang, Muhaoo Chen, Nanyun Peng, Kai-Wei Chang. Vulnerability of Large Language Models to Output Prefix Jailbreaks: Impact of Positions on Safety, *NAACL 2025*.
- [37] Lingrui Mei, Shenghua Liu, Yiwei Wang, Baolong Bi, Jiayi Mao, Xueqi Cheng. "Not Aligned" is Not "Malicious": Being Careful about Hallucinations of Large Language Models' Jailbreak, *COLING 2025*.
- [38] Cheng Wang, Yiwei Wang, Bryan Hooi, Yujun Cai, Nanyun Peng, Kai-Wei Chang. Con-ReCall: Detecting Pre-training Data in LLMs via Contrastive Decoding, *COLING 2025*.
- [39] Yihong Luo, Yuhang Chen, Siya Qiu, Yiwei Wang, Chen Zhang, Yan Zhou, Xiaochun Cao, Jing Tang. Fast Graph Sharpness-Aware Minimization for Enhancing and Accelerating Few-Shot Node Classification, *NeurIPS 2024*.
- [40] Bingxuan Li, Yiwei Wang, Tao Meng, Nanyun Peng, Kai-Wei Chang. Control Large Language Models via Divide-and-Conquer, *EMNLP 2024*.
- [41] Lingrui Mei, Shenghua Liu, Yiwei Wang, Baolong Bi, Xueqi Chen. SLANG: New Concept Comprehension of Large Language Models, *EMNLP 2024*.
- [42] Silin Meng, Yiwei Wang, Cheng-Fu Yang, Nanyun Peng, Kai-Wei Chang. LLM-A*: Large Language Model Enhanced Incremental Heuristic Search on Path Planning, *EMNLP 2024*.
- [43] Zhicheng Yang, Yinya Huang, Jing Xiong, Liang Feng, Xiaodan Liang, Yiwei Wang, Jing Tang. AlignedCoT: Prompting Large Language Models via Native-Speaking Demonstrations, *EMNLP 2024*.
- [44] Baolong Bi, Shenghua Liu, Yiwei Wang, Lingrui Mei, Xueqi Cheng. Lpnl: Scalable link prediction with large language models, *ACL 2024*.
- [45] Juncheng Liu, Bryan Hooi, Kenji Kawaguchi, Yiwei Wang, Chaosheng Dong, Xiaokui Xiao. Scalable and Effective Implicit Graph Neural Networks on Large Graphs, *ICLR 2024*.
- [46] Fei Wang, Wenjie Mo, Yiwei Wang, Wenxuan Zhou, Muhaoo Chen. A Causal View of Entity Bias in (Large) Language Models, *EMNLP 2023*.
- [47] Yiwei Wang, Yujun Cai, Muhaoo Chen, Yuxuan Liang, Bryan Hooi. Primacy Effect of ChatGPT, *EMNLP 2023*.
- [48] Yuxuan Liang, Yutong Xia, Songyu Ke, Yiwei Wang, Qingsong Wen, Junbo Zhang, Yu Zheng, Roger Zimmermann. Airformer: Predicting nationwide air quality in china with transformers, *AAAI 2023*.
- [49] Yiwei Wang, Bryan Hooi, Fei Wang, Yujun Cai, Yuxuan Liang, Wenxuan Zhou, Jing Tang, Manjuan Duan, Muhaoo Chen. How Fragile is Relation Extraction under Entity Replacements?, *CONLL 2023*.
- [50] Yiwei Wang, Bryan Hooi, Yozen Liu, Neil Shah. Graph explicit neural networks: Explicitly encoding graphs for efficient and accurate inference, *WSDM 2023*.
- [51] Juncheng Liu, Zequn Sun, Bryan Hooi, Yiwei Wang, Dayiheng Liu, Baosong Yang, Xiaokui Xiao, Muhaoo Chen. Dangling-Aware Entity Alignment with Mixed High-Order Proximities, *NAACL 2022*.

- [52] Yiwei Wang, Muhaoo Chen, Wenxuan Zhou, Yujun Cai, Yuxuan Liang, Bryan Hooi. GraphCache: Message Passing as Caching for Sentence-Level Relation Extraction, *NAACL 2022*.
- [53] Yiwei Wang, Muhaoo Chen, Wenxuan Zhou, Yujun Cai, Yuxuan Liang, Dayiheng Liu, Baosong Yang, Juncheng Liu, Bryan Hooi. Should We Rely on Entity Mentions for Relation Extraction? Debiasing Relation Extraction with Counterfactual Analysis, *NAACL 2022*.
- [54] Juncheng Liu, Yiwei Wang, Bryan Hooi, Renchi Yang, Xiaokui Xiao. Lscale: latent space clustering-based active learning for node classification, *ECML 2022*.
- [55] Yiwei Wang, Bryan Hooi, Yozen Liu, Tong Zhao, Zhichun Guo, Neil Shah. Flashlight: Scalable link prediction with effective decoders, *LOG 2022*.
- [56] Yuxuan Liang, Kun Ouyang, Yiwei Wang, Zheyi Pan, Yifang Yin, Hongyang Chen, Junbo Zhang, Yu Zheng, David S Rosenblum, Roger Zimmermann. Mixed-order relation-aware recurrent neural networks for spatio-temporal forecasting, *TKDE 2022*.
- [57] Juncheng Liu, Kenji Kawaguchi, Bryan Hooi, Yiwei Wang, Xiaokui Xiao. Eignn: Efficient infinite-depth graph neural networks, *NeurIPS 2021*.
- [58] Yiwei Wang, Yujun Cai, Yuxuan Liang, Henghui Ding, Changhu Wang, Siddharth Bhatia, Bryan Hooi. Adaptive data augmentation on temporal graphs, *NeurIPS 2021*.
- [59] Yujun Cai, Yiwei Wang, Yiheng Zhu, Tat-Jen Cham, Jianfei Cai, Junsong Yuan, Jun Liu, Chuanxia Zheng, Sijie Yan, Henghui Ding, others. A unified 3d human motion synthesis model via conditional variational auto-encoder, *ICCV 2021*.
- [60] Yiwei Wang, Yujun Cai, Yuxuan Liang, Henghui Ding, Changhu Wang, Bryan Hooi. Time-aware neighbor sampling for temporal graph networks, *2022 International Joint Conference on Neural Networks (IJCNN) 2021*.
- [61] Yuxuan Liang, Kun Ouyang, Yiwei Wang, Ye Liu, Junbo Zhang, Yu Zheng, David S Rosenblum. Revisiting convolutional neural networks for citywide crowd flow analytics, *ECML 2021*.
- [62] Yiwei Wang, Wei Wang, Yuxuan Liang, Yujun Cai, Bryan Hooi. Progressive supervision for node classification, *ECML 2021*.
- [63] Yuxuan Liang, Kun Ouyang, Hanshu Yan, Yiwei Wang, Zekun Tong, Roger Zimmermann. Modeling Trajectories with Neural Ordinary Differential Equations, *IJCAI 2021*.
- [64] Yuxuan Liang, Kun Ouyang, Junkai Sun, Yiwei Wang, Junbo Zhang, Yu Zheng, David Rosenblum, Roger Zimmermann. Fine-grained urban flow prediction, *WWW 2021*.
- [65] Yiwei Wang, Wei Wang, Yuxuan Liang, Yujun Cai, Bryan Hooi. Curgraph: Curriculum learning for graph classification, *WWW 2021*.
- [66] Yiwei Wang, Wei Wang, Yuxuan Liang, Yujun Cai, Bryan Hooi. Mixup for node and graph classification, *WWW 2021*.
- [67] Yuxuan Liang, Kun Ouyang, Yiwei Wang, David Samuel Rosenblum. Revisiting convolutional neural networks for urban flow analytics, *2020*.

- [68] Yujun Cai, Lin Huang, Yiwei Wang, Tat-Jen Cham, Jianfei Cai, Junsong Yuan, Jun Liu, Xu Yang, Yiheng Zhu, Xiaohui Shen, others. Learning progressive joint propagation for human motion prediction, *ECCV 2020*.
- [69] Yiwei Wang, Shenghua Liu, Minji Yoon, Hemank Lamba, Wei Wang, Christos Faloutsos, Bryan Hooi. Provably robust node classification via low-pass message passing, *ICDM 2020*.
- [70] Yiwei Wang, Wei Wang, Yujun Ca, Bryan Hooi, Beng Chin Ooi. Detecting implementation bugs in graph convolutional network based node classifiers, *ISSRE 2020*.
- [71] Yiwei Wang, Wei Wang, Yuxuan Liang, Yujun Cai, Juncheng Liu, Bryan Hooi. NodeAug: Semi-supervised node classification with data augmentation, *KDD 2020*.
- [72] Licheng Zhao, Yiwei Wang, Sandeep Kumar, Daniel P Palomar. Optimization algorithms for graph Laplacian estimation via ADMM and MM, *TSP 2019*.
- [73] Yiwei Wang, Mark James Carman, Yuan-Fang Li. Using knowledge graphs to explain entity co-occurrence in Twitter, *CIKM 2017*.

TEACHING

Instructor, New Course on Large Language Models	University of California, Merced, 2025
Teaching Assistant, Big Data Systems for Data Science	National University of Singapore, 2021
Teaching Assistant, Big Data Systems for Data Science	National University of Singapore, 2021
Teaching Assistant, Programming Methodology	National University of Singapore, 2021
Teaching Assistant, Knowledge Discovery and Data Mining and 2021	National University of Singapore, 2020
Teaching Assistant, Parallel Computing	National University of Singapore, 2020
Teaching Assistant, Signal Processing and Communications	HKUST, 2018

TUTORIAL

Towards Comprehensive Reasoning in Vision-Language Models	ICCV 2025
---	-----------

SERVING AS JOURNAL ASSOCIATE EDITOR (AE)

IEEE Transactions on Industrial Informatics (IEEE TII)	2025
Pattern Recognition	2025
Neurocomputing	2025

SERVING AS (SENIOR) AREA CHAIR

Association for the Advancement of Artificial Intelligence (AAAI)	2025
Conference on Neural Information Processing Systems (NeurIPS)	2025
International Conference on Learning Representations (ICLR)	2025

International Conference on Machine Learning (ICML)	2025
Association for Computational Linguistics Rolling Review (ACL ARR)	2024, 2025
Conference on Empirical Methods in Natural Language Processing (EMNLP)	2025
Conference on Language Modeling (COLM)	2025
International Joint Conference on Artificial Intelligence (IJCAI)	2025
IEEE International Conference on Multimedia and Expo (ICME)	2025
International Joint Conference on Neural Networks (IJCNN)	2025

SERVING AS REVIEWER

Conference on Empirical Methods in Natural Language Processing (EMNLP)	2023, 2024
Annual Meeting of the Association for Computational Linguistics (ACL)	2023, 2024
North American Chapter of the Association for Computational Linguistics (NAACL)	2023, 2024
ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD)	2023, 2024
International Conference on Artificial Intelligence and Statistics (AISTATS)	2025
Learning on Graphs Conference (LOG)	2022, 2023, 2024
The Web Conference (WWW)	2023, 2024, 2025
International Conference on Learning Representations (ICLR)	2023, 2024
Conference on Neural Information Processing Systems (NeurIPS)	2022, 2023, 2024
The AAAI Conference on Artificial Intelligence	2021, 2022, 2023, 2024
IEEE/CVF Conference on Computer Vision and Pattern Recognition	2021, 2022, 2023, 2024
International Conference on Computer Vision	2021, 2022, 2023, 2024
European Conference on Computer Vision	2022, 2023, 2024
International Joint Conference on Artificial Intelligence	2022, 2023, 2024
IEEE International Conference on Multimedia and Expo	2022, 2023, 2024
European Conference on Machine Learning and Principles and Practice of Knowledge Discovery in Databases (ECML-PKDD)	2020
International Joint Conference on Neural Networks	2022
IEEE Transactions on Pattern Analysis and Machine Intelligence (TPAMI)	
Neurocomputing	
Pattern Recognition	
IEEE Transactions on Knowledge and Data Engineering (TKDE)	

IET Image Processing

IEEE Transactions on Big Data

Computer Science Review

Cybernetics and Systems

Transactions on Systems, Man and Cybernetics:Systems