

王艺炜

电子邮箱: wangyw.evan@gmail.com 电话号码: +86 199-5526-8139 个人主页: <https://wangywust.github.io/>

工作经历

- **加州大学默塞德分校, 助理教授** 2025 年 1 月 – 至今
计算机科学系, 研究与教学方向为自然语言处理与视觉语言模型。
- **加州大学洛杉矶分校, 博士后研究员** 2023 年 9 月 – 2024 年 12 月
导师: Kai-Wei Chang 教授、Nanyun Peng 教授。开展大型语言模型与知识编辑相关研究。
- **亚马逊公司, 应用科学家** 2022 年 9 月 – 2023 年 9 月

教育背景

- **新加坡国立大学, 计算机科学博士 (Ph.D.)** 2019 年 7 月 – 2023 年 3 月
- **香港科技大学, 电子与计算机工程研究型硕士 (MPhil)** 2017 年 8 月 – 2019 年 6 月
- **东南大学, 信息工程学士** 2013 年 8 月 – 2017 年 6 月

荣誉与奖项

- 新加坡国立大学 **研究生院杰出研究奖** (全校博士生最高荣誉, 每年 3-4 名) 2022 年 7 月
- 新加坡数据科学学会 **杰出博士奖学金** (每年 10 名) 2021 年 3 月
- 新加坡国立大学计算学院 **研究成就奖** 2020 年 9 月
- 中华人民共和国教育部 **国家奖学金** (前 3%) 2015 年 10 月

教学经历

- 加州大学默塞德分校, 大语言模型课程主讲教师 2025 年
- 新加坡国立大学, 大数据系统课程助教 2021 年
- 新加坡国立大学, 编程方法课程助教 2021 年
- 新加坡国立大学, 知识发现与数据挖掘课程助教 2020, 2021 年
- 新加坡国立大学, 并行计算课程助教 2020 年
- 香港科技大学, 信号处理与通信课程助教 2018 年

学术期刊副编辑 (Journal Associate Editor)

- Neurocomputing 2025 年
- Pattern Recognition 2025 年

学术会议 (高级) 领域主席 ((Senior) Area Chair)

- NeurIPS 国际会议 2025 年
- ICLR 国际会议 2025 年
- ICML 国际会议 2025 年
- ACL Rolling Review 2024, 2025 年
- COLM 语言建模会议 2025 年
- IJCAI 人工智能联合会议 2025 年

- | | |
|--------------------|--------|
| • ICME 国际多媒体会议 | 2025 年 |
| • IJCNN 国际神经网络联合会议 | 2025 年 |

会议与期刊审稿服务

-
- EMNLP, ACL, NAACL, KDD, ICLR, NeurIPS, AAAI, CVPR, ICCV, ECCV, IJCAI, ICME, ECML-PKDD 等国际会议审稿人 (2021-2025)
 - IEEE TPAMI、TKDE, Pattern Recognition, Neurocomputing, IEEE Big Data 等国际期刊审稿人

研究论文发表

-
- [1] Xuan Shen, Brian Wingrenroth, Zichao Wang, Jason Kuen, Wanrong Zhu, Ruiyi Zhang, Yiwei Wang, Lichun Ma, Anqi Liu, Hongfu Liu, Tong Sun, Kevin S. Hawkins, Kate Tasker, G. Caleb Alexander, Jiuxiang Gu. OIDA-QA: A Multimodal Benchmark for Analyzing the Opioid Industry Documents Archive, *AAAI 2026*. (CCF-A)
 - [2] Yifan Lu, Pengfei Jiao, Xuan Guo, Ziyun Zou, Yiwei Wang, Mengzhou Gao, Huaming Wu, Imran Razzak. Towards Robust Heterogeneous Graph Explanations under Structural Perturbations, *WWW 2026*. (CCF-A)
 - [3] Juncheng Liu, Chenghao Liu, Gerald Woo, Yiwei Wang, Bryan Hooi, Caiming Xiong, Doyen Sahoo. UniTST: Effectively Modeling Inter-Series and Intra-Series Dependencies for Multivariate Time Series Forecasting, *TMLR 2025*.
 - [4] Jiayu Yao, Shenghua Liu, Yiwei Wang, Lingrui Mei, Baolong Bi, Yuyao Ge, Zhecheng Li, Xueqi Cheng. Who is in the Spotlight: The Hidden Bias Undermining Multimodal Retrieval-Augmented Generation, *EMNLP 2025*.
 - [5] Zhen Xiong, Yujun Cai, Zhecheng Li, Yiwei Wang. Mapping the Minds of LLMs: A Graph-Based Analysis of Reasoning LLMs, *EMNLP 2025*.
 - [6] Shuyang Hao, Yiwei Wang, Bryan Hooi, Jun Liu, Muhan Chen, Zi Huang, Yujun Cai. Making every step effective: Jailbreaking large vision-language models through hierarchical kv equalization, *EMNLP 2025*.
 - [7] Haonan Ge, Yiwei Wang, Ming-Hsuan Yang, Yujun Cai. MRFD: Multi-Region Fusion Decoding with Self-Consistency for Mitigating Hallucinations in LVLMs, *EMNLP 2025*.
 - [8] Hang Wu, Hongkai Chen, Yujun Cai, Chang Liu, Qingwen Ye, Ming-Hsuan Yang, Yiwei Wang. DiMo-GUI: Advancing Test-time Scaling in GUI Grounding via Modality-Aware Visual Reasoning, *EMNLP 2025*.
 - [9] Wen Tao, Yiwei Wang. How to Make Large Language Models Generate 100
 - [10] Sifan Li, Yujun Cai, Yiwei Wang. SemVink: Advancing VLMs' Semantic Understanding of Optical Illusions via Visual Global Thinking, *EMNLP 2025*.
 - [11] Xingjian Tao, Yiwei Wang, Yujun Cai, Zhicheng Yang, Jing Tang. Understanding GUI Agent Localization Biases through Logit Sharpness, *EMNLP 2025*.
 - [12] Zhecheng Li, Guoxian Song, Yujun Cai, Zhen Xiong, Junsong Yuan, Yiwei Wang. Texture or Semantics? Vision-Language Models Get Lost in Font Recognition, *COLM 2025*.
 - [13] Zhaochen Wang, Bryan Hooi, Yiwei Wang, Ming-Hsuan Yang, Zi Huang, Yujun Cai. Text Speaks Louder than Vision: ASCII Art Reveals Textual Biases in Vision-Language Models, *COLM 2025*.
 - [14] Chunxue Xu, Yiwei Wang, Bryan Hooi, Yujun Cai, Songze Li. How does Watermarking Affect Visual Language Models in Document Understanding?, *COLM 2025*.
 - [15] Baolong Bi, Shaohan Huang, Yiwei Wang, Tianchi Yang, Zihan Zhang, Haizhen Huang, Lingrui Mei, Junfeng Fang, Zehao Li, Furu Wei, Weiwei Deng, Feng Sun, Qi Zhang, Shenghua Liu. Context-DPO: Aligning Language Models for Context-Faithfulness, *ACL 2025*. (CCF-A)

- [16] Baolong Bi, Shenghua Liu, Lingrui Mei, Yiwei Wang, Junfeng Fang, Pengliang Ji, Xueqi Cheng. Decoding by Contrasting Knowledge: Enhancing Large Language Model Confidence on Edited Facts, *ACL 2025*. (CCF-A)
- [17] Yuyao Ge, Shenghua Liu, Baolong Bi, Yiwei Wang, Lingrui Mei, Wenjie Feng, Lizhe Chen, Xueqi Cheng. Can Graph Descriptive Order Affect Solving Graph Problems with LLMs?, *ACL 2025*. (CCF-A)
- [18] Zhecheng Li, Yiwei Wang, Bryan Hooi, Yujun Cai, Nanyun Peng, Kai-Wei Chang. DRS: Deep Question Reformulation With Structured Output, *ACL 2025*. (CCF-A)
- [19] Bingxuan Li, Yiwei Wang, Jiuxiang Gu, Kai-Wei Chang, Nanyun Peng. METAL: A Multi-Agent Framework for Chart Generation with Test-Time Scaling, *ACL 2025*. (CCF-A)
- [20] Zhecheng Li, Yiwei Wang, Bryan Hooi, Yujun Cai, Zhen Xiong, Nanyun Peng, Kai-Wei Chang. Vulnerability of LLMs to Vertically Aligned Text Manipulations, *ACL 2025*. (CCF-A)
- [21] Yurong Wu, Fangwen Mu, QiuHong Zhang, Jinjing Zhao, Xinrun Xu, Lingrui Mei, Yang Wu, Lin Shi, Junjie Wang, Zhiming Ding, Yiwei Wang. EvoStealer: Differential Evolution for Prompt Template Stealing Against Text-to-Image Synthesis, *ACL 2025*. (CCF-A)
- [22] Rongzhi Zhu, Xiangyu Liu, Zequn Sun, Yiwei Wang, Wei Hu. Mitigating Lost-in-Retrieval Problems in Retrieval Augmented Multi-Hop Question Answering, *ACL 2025*. (CCF-A)
- [23] Songlin Zhai, Yuan Meng, Yongrui Chen, Yiwei Wang, Guilin Qi. Peripheral Memory for LLMs: Integration of Sequential Memory Banks with Adaptive Querying, *ICML 2025*. (CCF-A)
- [24] Baolong Bi, Shenghua Liu, Yiwei Wang, Lingrui Mei, Xueqi Cheng. Is Factuality Enhancement a Free Lunch For LLMs? Better Factuality Can Lead to Worse Context-Faithfulness, *ICLR 2025*.
- [25] Zhicheng Yang, Yiwei Wang, Yinya Huang, Zhijiang Guo, Wei Shi, Xiongwei Han, Liang Feng, Linqi Song, Xiaodan Liang, Jing Tang. OptiBench meets ReSocratic: Measure and improve LLMs for optimization modeling, *ICLR 2025*.
- [26] Shaochen Zhong, others. MQuAKE-Remastered: Multi-Hop Knowledge Editing Can Only Be Advanced With Reliable Evaluations, *ICLR 2025*.
- [27] Cheng Wang, Yiwei Wang, Yujun Cai, Bryan Hooi. Tricking Retrievers with Influential Tokens: An Efficient Black-Box Corpus Poisoning Attack, *NAACL 2025*.
- [28] Yiwei Wang, Muhan Chen, Nanyun Peng, Kai-Wei Chang. Vulnerability of Large Language Models to Output Prefix Jailbreaks: Impact of Positions on Safety, *NAACL 2025*.
- [29] Lingrui Mei, Shenghua Liu, Yiwei Wang, Baolong Bi, Jiayi Mao, Xueqi Cheng. "Not Aligned" is Not "Malicious": Being Careful about Hallucinations of Large Language Models' Jailbreak, *COLING 2025*.
- [30] Cheng Wang, Yiwei Wang, Bryan Hooi, Yujun Cai, Nanyun Peng, Kai-Wei Chang. Con-ReCall: Detecting Pre-training Data in LLMs via Contrastive Decoding, *COLING 2025*.
- [31] Yihong Luo, Yuhua Chen, Siya Qiu, Yiwei Wang, Chen Zhang, Yan Zhou, Xiaochun Cao, Jing Tang. Fast Graph Sharpness-Aware Minimization for Enhancing and Accelerating Few-Shot Node Classification, *NeurIPS 2024*. (CCF-A)
- [32] Bingxuan Li, Yiwei Wang, Tao Meng, Nanyun Peng, Kai-Wei Chang. Control Large Language Models via Divide-and-Conquer, *EMNLP 2024*.
- [33] Lingrui Mei, Shenghua Liu, Yiwei Wang, Baolong Bi, Xueqi Chen. SLANG: New Concept Comprehension of Large Language Models, *EMNLP 2024*.
- [34] Silin Meng, Yiwei Wang, Cheng-Fu Yang, Nanyun Peng, Kai-Wei Chang. LLM-A*: Large Language Model Enhanced Incremental Heuristic Search on Path Planning, *EMNLP 2024*.

- [35] Zhicheng Yang, Yinya Huang, Jing Xiong, Liang Feng, Xiaodan Liang, Yiwei Wang, Jing Tang. AlignedCoT: Prompting Large Language Models via Native-Speaking Demonstrations, *EMNLP 2024*.
- [36] Baolong Bi, Shenghua Liu, Yiwei Wang, Lingrui Mei, Xueqi Cheng. Lpnl: Scalable link prediction with large language models, *ACL 2024*. (CCF-A)
- [37] Juncheng Liu, Bryan Hooi, Kenji Kawaguchi, Yiwei Wang, Chaosheng Dong, Xiaokui Xiao. Scalable and Effective Implicit Graph Neural Networks on Large Graphs, *ICLR 2024*.
- [38] Fei Wang, Wenjie Mo, Yiwei Wang, Wenxuan Zhou, Muhaao Chen. A Causal View of Entity Bias in (Large) Language Models, *EMNLP 2023*.
- [39] Yiwei Wang, Yujun Cai, Muhaao Chen, Yuxuan Liang, Bryan Hooi. Primacy Effect of ChatGPT, *EMNLP 2023*.
- [40] Yuxuan Liang, Yutong Xia, Songyu Ke, Yiwei Wang, Qingsong Wen, Junbo Zhang, Yu Zheng, Roger Zimmermann. Airformer: Predicting nationwide air quality in china with transformers, *AAAI 2023*. (CCF-A)
- [41] Yiwei Wang, Bryan Hooi, Fei Wang, Yujun Cai, Yuxuan Liang, Wenxuan Zhou, Jing Tang, Manjuan Duan, Muhaao Chen. How Fragile is Relation Extraction under Entity Replacements?, *CONLL 2023*.
- [42] Yiwei Wang, Bryan Hooi, Yozen Liu, Neil Shah. Graph explicit neural networks: Explicitly encoding graphs for efficient and accurate inference, *WSDM 2023*.
- [43] Juncheng Liu, Zequn Sun, Bryan Hooi, Yiwei Wang, Dayiheng Liu, Baosong Yang, Xiaokui Xiao, Muhaao Chen. Dangling-Aware Entity Alignment with Mixed High-Order Proximities, *NAACL 2022*.
- [44] Yiwei Wang, Muhaao Chen, Wenxuan Zhou, Yujun Cai, Yuxuan Liang, Bryan Hooi. GraphCache: Message Passing as Caching for Sentence-Level Relation Extraction, *NAACL 2022*.
- [45] Yiwei Wang, Muhaao Chen, Wenxuan Zhou, Yujun Cai, Yuxuan Liang, Dayiheng Liu, Baosong Yang, Juncheng Liu, Bryan Hooi. Should We Rely on Entity Mentions for Relation Extraction? Debiasing Relation Extraction with Counterfactual Analysis, *NAACL 2022*.
- [46] Juncheng Liu, Yiwei Wang, Bryan Hooi, Renchi Yang, Xiaokui Xiao. Lscale: latent space clustering-based active learning for node classification, *ECML 2022*.
- [47] Yiwei Wang, Bryan Hooi, Yozen Liu, Tong Zhao, Zhichun Guo, Neil Shah. Flashlight: Scalable link prediction with effective decoders, *LOG 2022*.
- [48] Yuxuan Liang, Kun Ouyang, Yiwei Wang, Zheyi Pan, Yifang Yin, Hongyang Chen, Junbo Zhang, Yu Zheng, David S Rosenblum, Roger Zimmermann. Mixed-order relation-aware recurrent neural networks for spatio-temporal forecasting, *TKDE 2022*. (CCF-A)
- [49] Juncheng Liu, Kenji Kawaguchi, Bryan Hooi, Yiwei Wang, Xiaokui Xiao. Eignn: Efficient infinite-depth graph neural networks, *NeurIPS 2021*. (CCF-A)
- [50] Yiwei Wang, Yujun Cai, Yuxuan Liang, Henghui Ding, Changhu Wang, Siddharth Bhatia, Bryan Hooi. Adaptive data augmentation on temporal graphs, *NeurIPS 2021*. (CCF-A)
- [51] Yujun Cai, Yiwei Wang, Yiheng Zhu, Tat-Jen Cham, Jianfei Cai, Junsong Yuan, Jun Liu, Chuanxia Zheng, Sijie Yan, Henghui Ding, others. A unified 3d human motion synthesis model via conditional variational auto-encoder, *ICCV 2021*. (CCF-A)
- [52] Yiwei Wang, Yujun Cai, Yuxuan Liang, Henghui Ding, Changhu Wang, Bryan Hooi. Time-aware neighbor sampling for temporal graph networks, *2022 International Joint Conference on Neural Networks (IJCNN) 2021*.
- [53] Yuxuan Liang, Kun Ouyang, Yiwei Wang, Ye Liu, Junbo Zhang, Yu Zheng, David S Rosenblum. Revisiting convolutional neural networks for citywide crowd flow analytics, *ECML 2021*.
- [54] Yiwei Wang, Wei Wang, Yuxuan Liang, Yujun Cai, Bryan Hooi. Progressive supervision for node classification,

ECML 2021.

- [55] Yuxuan Liang, Kun Ouyang, Hanshu Yan, Yiwei Wang, Zekun Tong, Roger Zimmermann. Modeling Trajectories with Neural Ordinary Differential Equations, *IJCAI 2021*. (CCF-A)
- [56] Yuxuan Liang, Kun Ouyang, Junkai Sun, Yiwei Wang, Junbo Zhang, Yu Zheng, David Rosenblum, Roger Zimmermann. Fine-grained urban flow prediction, *WWW 2021*. (CCF-A)
- [57] Yiwei Wang, Wei Wang, Yuxuan Liang, Yujun Cai, Bryan Hooi. Curgraph: Curriculum learning for graph classification, *WWW 2021*. (CCF-A)
- [58] Yiwei Wang, Wei Wang, Yuxuan Liang, Yujun Cai, Bryan Hooi. Mixup for node and graph classification, *WWW 2021*. (CCF-A)
- [59] Yuxuan Liang, Kun Ouyang, Yiwei Wang, David Samuel Rosenblum. Revisiting convolutional neural networks for urban flow analytics, 2020.
- [60] Yujun Cai, Lin Huang, Yiwei Wang, Tat-Jen Cham, Jianfei Cai, Junsong Yuan, Jun Liu, Xu Yang, Yiheng Zhu, Xiaohui Shen, others. Learning progressive joint propagation for human motion prediction, *ECCV 2020*. (CCF-A)
- [61] Yiwei Wang, Shenghua Liu, Minji Yoon, Hemank Lamba, Wei Wang, Christos Faloutsos, Bryan Hooi. Provably robust node classification via low-pass message passing, *ICDM 2020*.
- [62] Yiwei Wang, Wei Wang, Yujun Ca, Bryan Hooi, Beng Chin Ooi. Detecting implementation bugs in graph convolutional network based node classifiers, *ISSRE 2020*.
- [63] Yiwei Wang, Wei Wang, Yuxuan Liang, Yujun Cai, Juncheng Liu, Bryan Hooi. NodeAug: Semi-supervised node classification with data augmentation, *KDD 2020*. (CCF-A)
- [64] Licheng Zhao, Yiwei Wang, Sandeep Kumar, Daniel P Palomar. Optimization algorithms for graph Laplacian estimation via ADMM and MM, *TSP 2019*. (CCF-A)
- [65] Yiwei Wang, Mark James Carman, Yuan-Fang Li. Using knowledge graphs to explain entity co-occurrence in Twitter, *CIKM 2017*.